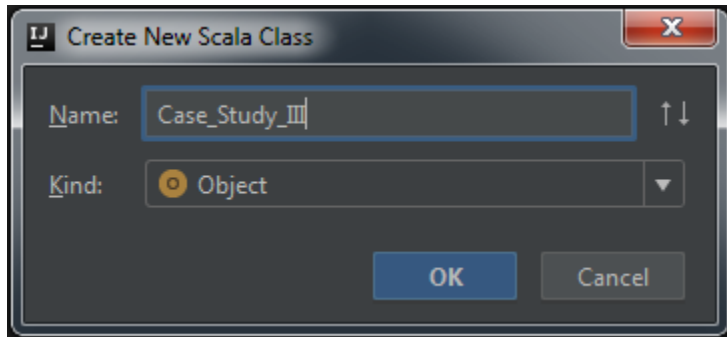


Case Study III for Session 20 – Spark SQL

- Prachi Mohite

In this Assignment we will be using IDEA IntelliJ to Complete the given Task

1. Created new Project and added scala object named as Case_Study_III as below



2. To add the required dependencies we have created scala sbt project in IDEA and added library dependency from maven repository as below

```
name := "Project1"

version := "0.1"

scalaVersion := "2.11.7"
libraryDependencies += "org.apache.spark" %% "spark-core" % "2.1.0"
libraryDependencies += "org.apache.spark" %% "spark-sql" % "2.1.0" % "provided"
```

3. Added main function and created the spark object as below


```
def main(args: Array[String]): Unit = {

    println("hey scala")


    //Create spark object
    val spark = SparkSession
        .builder()
        .master( master = "local")
        .appName( name = "Spark SQL basic example")
        .config("spark.some.config.option", "some-value")
        .getOrCreate()

    println("Spark Session Object created")
}
```

4. We will be using below dataset for this assignment
 - a. Building CSV
 - b. Columns are – BuildingID, BuildingMgr, BuildingAge, HVACproduct, country



- c. HVAC.csv
- d. Columns are – Date, Time , TargetTemp, ActualTemp, System, SystemAge, BuildingID



Task 1

Load HVAC.csv into temporary table.

Solution Approach –

In this case we have to first create a Dataframe and then register that data frame as a temporary table. This can be achieved by two approaches as below

- **Inferring the Schema using Reflection:** This method uses reflection to generate the schema of an RDD that contains specific types of objects.
 - In this case we have create the case class which will reflect the schema of dataframe / temp table which should hold values from the file.
 - Once case class is created, we have to read the text file and create a RDD based on spit operator specified for that file data. Apply transformations on RDD and get the required Dataframe created.
 - Create a temporary table from the Data frame.

Note : Above approach we have applied in Assignment 20 , 21.

- **Programmatically Specifying the Schema:** The second method for creating DataFrame is through programmatic interface that allows you to construct a schema and then apply it to an existing RDD
 - In this method we have to create a manual schema which should matches to our requirement and data of file to be loaded.
 - Create a Data frame using createDataFrame method on RowRDD or while loading only apply schema of load method

Note : In this case study we will be using this Approach

1. Create a manual schema which matches data in both the files building.csv and HVAC.csv
 - a. While creating the schema we have used Struct type.
 - b. StructType is a built-in data type used for Schema definition in Spark SQL, to represent a collection of StructFields that together define a schema or its part.

<schema-name> = new

structType<array_of_columns><Struct_field>(<column_name>,<data_type_of_column>,<nullable_or_not_nullable(true/false)>)

```
2. //Schema for HVAC.csv
val Manual_schema_HVAC = new StructType(Array(new
StructField("Date", StringType,
true),
new StructField("Time", StringType, false),
new StructField("TargetTemp", LongType, true),
new StructField("ActualTemp", LongType, false),
new StructField("System", LongType, false),
new StructField("SystemAge", LongType, false),
new StructField("BuildingID", LongType, false)))

//Schema for building.csv
val Manual_schema_Building = new StructType(Array(new
StructField("BuildingID",
LongType, true),
new StructField("BuildingMgr", StringType, false),
new StructField("BuildingAge", LongType, true),
new StructField("HVACproduct", StringType, false),
new StructField("Country", StringType, false)))
```

3. Load the the CSV files from local system to Spark as below
 - a. Spark SQL provides inbuilt support for only 3 types of data sources
 - i. Parquet (This is default)
 - ii. JSON
 - iii. JDBC
 - iv. For CSV we have to use another library spark-csv this was prior to Spark 2.0. As of Spark version 2.0 and up, spark-csv is part of core Spark functionality and doesn't require a separate library.

Below code is how to read and load the csv file and create a dataframe

We are using the CSV file read format, this provides various options of which we have used a few of them, which are as follows:

- option to remove the header from the input file.
- given the manual schema that we have created in the previous step
- the path where the CSV file is saved in the local file system.

```
//Reading the HVAC csv file
val HVAC = spark.read.format("CSV")
    .option("header",true)
    .schema(Manual_schema_HVAC)
    .load("E:\\Prachi IMP\\Hadoop\\Case Studies - Assignment\\Case Study
III\\HVAC.csv")

//Displaying the dataframe contents
HVAC.show()
```

Output of the show command

```
16/03/20 10:00:30 INFO CodeGenerator: Code generated in 22.177447 ms
+-----+-----+-----+-----+-----+-----+-----+
| Date | Time | TargetTemp | ActualTemp | System | SystemAge | BuildingID |
+-----+-----+-----+-----+-----+-----+-----+
| 6/1/13 | 0:00:01 | 66 | 58 | 13 | 20 | 4 |
| 6/2/13 | 1:00:01 | 69 | 68 | 3 | 20 | 17 |
| 6/3/13 | 2:00:01 | 70 | 73 | 17 | 20 | 18 |
| 6/4/13 | 3:00:01 | 67 | 63 | 2 | 23 | 15 |
| 6/5/13 | 4:00:01 | 68 | 74 | 16 | 9 | 3 |
| 6/6/13 | 5:00:01 | 67 | 56 | 13 | 28 | 4 |
| 6/7/13 | 6:00:01 | 70 | 58 | 12 | 24 | 2 |
| 6/8/13 | 7:00:01 | 70 | 73 | 20 | 26 | 16 |
| 6/9/13 | 8:00:01 | 66 | 69 | 16 | 9 | 9 |
| 6/10/13 | 9:00:01 | 65 | 57 | 6 | 5 | 12 |
| 6/11/13 | 10:00:01 | 67 | 70 | 10 | 17 | 15 |
| 6/12/13 | 11:00:01 | 69 | 62 | 2 | 11 | 7 |
| 6/13/13 | 12:00:01 | 69 | 73 | 14 | 2 | 15 |
| 6/14/13 | 13:00:01 | 65 | 61 | 3 | 2 | 6 |
| 6/15/13 | 14:00:01 | 67 | 59 | 19 | 22 | 20 |
| 6/16/13 | 15:00:01 | 65 | 56 | 19 | 11 | 8 |
| 6/17/13 | 16:00:01 | 67 | 57 | 15 | 7 | 6 |
| 6/18/13 | 17:00:01 | 66 | 57 | 12 | 5 | 13 |
| 6/19/13 | 18:00:01 | 69 | 58 | 8 | 22 | 4 |
| 6/20/13 | 19:00:01 | 67 | 55 | 17 | 5 | 7 |
+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows
```

4. Registering the temporary table with registerTempTable command

- a. registerTempTable() creates an in-memory table that is scoped to the cluster in which it was created. The data is stored using Hive's highly-optimized, in-memory columnar format.
- b. This is important for dashboards as dashboards running in a different cluster (ie. the single Dashboard Cluster) will not have access to the temp tables registered in another cluster.
- c. Re-registering a temp table of the same name (using overwrite=true) but with new data causes an atomic memory pointer switch so the new data is seamlessly updated and immediately accessible for querying (ie. from a Dashboard).

```
//Registering the temporary table
HVAC.registerTempTable("HVAC_table")
println("HVAC table registered!")
```

Output

```
18/05/28 16:06:56 INFO BlockM
18/05/28 16:06:56 INFO BlockM
HVAC table registered!
18/05/28 16:06:56 INFO FileSo
18/05/28 16:06:56 INFO FileSo
18/05/28 16:06:56 INFO FileSo
```

5. Same above steps are followed to create the dataframe for building.csv file and create a temp table

Code

```
//Reading the building.csv file and creating the temp table
val buildings = spark.read.format("CSV")
    .option("header", true)
    .schema(Manual_schema_Building)
    .load("E:\\Prachi IMP\\Hadoop\\Case Studies - Assignment\\Case Study
III\\building.csv")
buildings.show()
buildings.registerTempTable("building_table")
println("buildings table registered!")
```

Output

```
18/05/28 16:06:56 INFO SparkSQLParser: Parsing Command: select *, IF((TargetTemp-ActualTemp)< -5 , '1',0)) as tempchange from HVAC_table
```

BuildingID	BuildingMgr	BuildingAge	HVACproduct	Country
1	M1	25	AC1000	USA
2	M2	27	FN39TG	France
3	M3	28	JDNS77	Brazil
4	M4	17	GG1919	Finland
5	M5	3	ACMAX22	Hong Kong
6	M6	9	AC1000	Singapore
7	M7	13	FN39TG	South Africa
8	M8	25	JDNS77	Australia
9	M9	11	GG1919	Mexico
10	M10	23	ACMAX22	China
11	M11	14	AC1000	Belgium
12	M12	26	FN39TG	Finland
13	M13	25	JDNS77	Saudi Arabia
14	M14	17	GG1919	Germany
15	M15	19	ACMAX22	Israel
16	M16	23	AC1000	Turkey
17	M17	11	FN39TG	Egypt
18	M18	25	JDNS77	Indonesia
19	M19	14	GG1919	Canada
20	M20	19	ACMAX22	Argentina

```
buildings table registered!
18/05/28 16:06:56 INFO FileSourceS
18/05/28 16:06:56 INFO FileSourceS
```

Task 1.2

Add a new column, tempchange - set to 1, if there is a change of greater than +/-5 between actual and target temperature

Solution Approach –

AS we have Temp Table and Dataframe created, a new column can be added with below two approaches

Approach 1: Using SQL queries on Temp Table

```
//Approach 1: Adding new column to temp table using SQL Queries
val filterHVAC = spark.sql(
  """select *, IF((TargetTemp-ActualTemp)> 5 , '1',
    |IF((TargetTemp-ActualTemp)< -5 , '1',0)) as Temp_change_diff from
HVAC_table""".stripMargin)
filterHVAC.show()
```

Output

18/05/28 16:06:56 INFO CodeGenerator: Code generated in 14.534886 ms

Date	Time	TargetTemp	ActualTemp	System	SystemAge	BuildingID	tempchange
6/1/13	0:00:01	66	58	13	20	4	1
6/2/13	1:00:01	69	68	3	20	17	0
6/3/13	2:00:01	70	73	17	20	18	0
6/4/13	3:00:01	67	63	2	23	15	0
6/5/13	4:00:01	68	74	16	9	3	1
6/6/13	5:00:01	67	56	13	28	4	1
6/7/13	6:00:01	70	58	12	24	2	1
6/8/13	7:00:01	70	73	20	26	16	0
6/9/13	8:00:01	66	69	16	9	9	0
6/10/13	9:00:01	65	57	6	5	12	1
6/11/13	10:00:01	67	70	10	17	15	0
6/12/13	11:00:01	69	62	2	11	7	1
6/13/13	12:00:01	69	73	14	2	15	0
6/14/13	13:00:01	65	61	3	2	6	0
6/15/13	14:00:01	67	59	19	22	20	1
6/16/13	15:00:01	65	56	19	11	8	1
6/17/13	16:00:01	67	57	15	7	6	1
6/18/13	17:00:01	66	57	12	5	13	1
6/19/13	18:00:01	69	58	8	22	4	1
6/20/13	19:00:01	67	55	17	5	7	1

Approach 2: Using Dataframe and 'withcoulmn' transformations to get new column which is added conditionally

```
//Approach 2: Using Dataframe and 'withCoulmn' Transformation
val filterHVAC2 = HVAC.withColumn("tempChange",when(
  ($"TargetTemp"-$"ActualTemp")> 5,1).otherwise(0)).toDF()
filterHVAC2.show()
```

Output


```

18/05/28 16:06:56 INFO DAGScheduler: Job 2 finished: show at Case_Study_III.scala:67, took 0.0266
18/05/28 16:06:56 INFO CodeGenerator: Code generated in 14.534886 ms
+-----+-----+-----+-----+-----+-----+-----+-----+
| Date | Time | TargetTemp | ActualTemp | System | SystemAge | BuildingID | tempchange |
+-----+-----+-----+-----+-----+-----+-----+-----+
| 6/1/13 | 0:00:01 | 66 | 58 | 13 | 20 | 4 | 1 |
| 6/2/13 | 1:00:01 | 69 | 68 | 3 | 20 | 17 | 0 |
| 6/3/13 | 2:00:01 | 70 | 73 | 17 | 20 | 18 | 0 |
| 6/4/13 | 3:00:01 | 67 | 63 | 2 | 23 | 15 | 0 |
| 6/5/13 | 4:00:01 | 68 | 74 | 16 | 9 | 3 | 1 |
| 6/6/13 | 5:00:01 | 67 | 56 | 13 | 28 | 4 | 1 |
| 6/7/13 | 6:00:01 | 70 | 58 | 12 | 24 | 2 | 1 |
| 6/8/13 | 7:00:01 | 70 | 73 | 20 | 26 | 16 | 0 |
| 6/9/13 | 8:00:01 | 66 | 69 | 16 | 9 | 9 | 0 |
| 6/10/13 | 9:00:01 | 65 | 57 | 6 | 5 | 12 | 1 |
| 6/11/13 | 10:00:01 | 67 | 70 | 10 | 17 | 15 | 0 |
| 6/12/13 | 11:00:01 | 69 | 62 | 2 | 11 | 7 | 1 |
| 6/13/13 | 12:00:01 | 69 | 73 | 14 | 2 | 15 | 0 |
| 6/14/13 | 13:00:01 | 65 | 61 | 3 | 2 | 6 | 0 |
| 6/15/13 | 14:00:01 | 67 | 59 | 19 | 22 | 20 | 1 |
| 6/16/13 | 15:00:01 | 65 | 56 | 19 | 11 | 8 | 1 |
| 6/17/13 | 16:00:01 | 67 | 57 | 15 | 7 | 6 | 1 |
| 6/18/13 | 17:00:01 | 66 | 57 | 12 | 5 | 13 | 1 |
| 6/19/13 | 18:00:01 | 69 | 58 | 8 | 22 | 4 | 1 |
| 6/20/13 | 19:00:01 | 67 | 55 | 17 | 5 | 7 | 1 |
+-----+-----+-----+-----+-----+-----+-----+-----+
only showing top 20 rows

```

Note: Instead of writing the conditions for adding new column inline in SQL Query or in withColumn transformation, we can write a udf and register it and we can use it as well

Task 2 Load the building.csv file in temporary table

As mentioned above already loaded building.csv as well

Code

```

//Reading the building.csv file and creating the temp table
val buildings = spark.read.format("CSV")
    .option("header", true)
    .schema(Manual_schema_Building)
    .load("E:\\casestudies\\sensorcasestudy\\buliding.csv")
buildings.show()
buildings.registerTempTable("building_table")
println("buildings table registered!")

```

Output

```
IF((TargetTemp-ActualTemp)< -5 , '1',0)) as tempchange from HVAC_table
```

BuildingID	BuildingMgr	BuildingAge	HVACproduct	Country
1	M1	25	AC1000	USA
2	M2	27	FN39TG	France
3	M3	28	JDNS77	Brazil
4	M4	17	GG1919	Finland
5	M5	3	ACMAX22	Hong Kong
6	M6	9	AC1000	Singapore
7	M7	13	FN39TG	South Africa
8	M8	25	JDNS77	Australia
9	M9	11	GG1919	Mexico
10	M10	23	ACMAX22	China
11	M11	14	AC1000	Belgium
12	M12	26	FN39TG	Finland
13	M13	25	JDNS77	Saudi Arabia
14	M14	17	GG1919	Germany
15	M15	19	ACMAX22	Israel
16	M16	23	AC1000	Turkey
17	M17	11	FN39TG	Egypt
18	M18	25	JDNS77	Indonesia
19	M19	14	GG1919	Canada
20	M20	19	ACMAX22	Argentina

```
buildings table registered!
```

```
18/05/28 16:06:56 INFO FileSourceStrategy: Pruning directories with:
```

Task 3.1 Figure out the number of times, temperature has changed by 5 degrees or more for each country

Solution Approach –

- We need to join both the tables on key building ID
- We have already created a column 'tempchange' in task 1.2 , select tempchange where it is = 1 and group on country to get the count

Approach 1 : Using SQL Queries

```
//Approach 1: Using SQL Queries
filterHVAC2.registerTempTable("FilteredHVAC")
spark.sql("Select Country,Count(Country) as Count from FilteredHVAC
HVAC " +
"Join building_table BLD On HVAC.BuildingID=BLD.BuildingID where
tempchange=1 group by country").show()
```

Output

```
18/05/28 16:06:59 INFO CodeGenerator: C
+-----+-----+
|   Country|Count|
+-----+-----+
|  Singapore|  109|
|   Turkey|  117|
|  Germany|   96|
|   France|  100|
| Argentina|  114|
|  Belgium|  102|
|  Finland|  231|
|   China|  117|
| Hong Kong|  123|
|  Israel|  106|
|   USA|  109|
|  Mexico|  122|
| Indonesia|  128|
| Saudi Arabia| 122|
|   Canada|   98|
|   Brazil|  105|
| Australia|   97|
|   Egypt|  124|
| South Africa| 104|
+-----+-----+

18/05/28 16:06:59 INFO FileSourceStrate
18/05/28 16:06:59 INFO FileSourceStrate
```

Approach 2: Using Spark SQL Transformations → join , filter , groupby , count

```
//Approach 2: Using Spark Transformations
filterHVAC2.as("HVAC").join(buildings.as("BLD"), $"HVAC.BuildingID"=== $"BLD.BuildingID")
  .filter($"tempchange"===1).groupBy("Country").count().show()
```

Output

```
18/05/28 16:07:01 INFO DAGScheduler: ResultStage 23 (show a
18/05/28 16:07:01 INFO DAGScheduler: Job 15 finished: show a
+-----+-----+
| Country | count |
+-----+-----+
| Singapore | 109 |
| Turkey | 117 |
| Germany | 96 |
| France | 100 |
| Argentina | 114 |
| Belgium | 102 |
| Finland | 231 |
| China | 117 |
| Hong Kong | 123 |
| Israel | 106 |
| USA | 109 |
| Mexico | 122 |
| Indonesia | 128 |
| Saudi Arabia | 122 |
| Canada | 98 |
| Brazil | 105 |
| Australia | 97 |
| Egypt | 124 |
| South Africa | 104 |
+-----+-----+

18/05/28 16:07:01 INFO FileSourceStrategy: Pruning directori
18/05/28 16:07:01 INFO FileSourceStrategy: Post-Scan Filters
18/05/28 16:07:01 INFO FileSourceStrategy: Pruning directori
```

Approach 3:

We can also create a Dataframe by joining two tables and register the newly created table joining and perform the querying on the same.

```
//Approach 3: Creating temp table after joining two tables
val joinExpression = filterHVAC.col("BuildingID") ===
    buildings.toDF().col("BuildingID")
val HVACJOBUILD = filterHVAC.join(buildings,joinExpression)
HVACJOBUILD.show()
HVACJOBUILD.registerTempTable("HVACJBUILD")
```

```

val selective = spark.sql("""select tempchange, Country from
HVACJBUILD WHERE
                                tempchange = 1""").toDF()

// registering the nwly created table
selective.registerTempTable("newselective")

//Approach 3.1
spark.sql("""select Country, count(tempchange) from newselective
group by
Country""").show()

//Approach 3.2
selective.filter($"tempchange"===1).groupBy("Country").count().show()

```

Output

Approach 3.2

```

18/05/28 16:07:04 INFO SparkContext: Invoking
+-----+-----+
|   Country|count|
+-----+-----+
| Singapore|   230|
|   Turkey|   243|
|   Germany|   196|
|   France|   251|
| Argentina|   230|
|   Belgium|   199|
|   Finland|   473|
|   China|   241|
| Hong Kong|   248|
|   Israel|   232|
|   USA|   213|
|   Mexico|   228|
| Indonesia|   243|
| Saudi Arabia|   233|
|   Canada|   232|
|   Brazil|   226|
| Australia|   225|
|   Egypt|   236|
| South Africa|   237|
+-----+-----+

18/05/28 16:07:02 INFO DAGScheduler: Job 23 finished: show
+-----+-----+
|   Country|count(tempchange)|
+-----+-----+
| Singapore|               230|
|   Turkey|               243|
|   Germany|               196|
|   France|               251|
| Argentina|               230|
|   Belgium|               199|
|   Finland|               473|
|   China|               241|
| Hong Kong|               248|
|   Israel|               232|
|   USA|               213|
|   Mexico|               228|
| Indonesia|               243|
| Saudi Arabia|               233|
|   Canada|               232|
|   Brazil|               226|
| Australia|               225|
|   Egypt|               236|
| South Africa|               237|
+-----+-----+

18/05/28 16:07:04 INFO SparkUI: Stopped Spark
18/05/28 16:07:04 INFO MapOutputTrackerMaster
18/05/28 16:07:04 INFO MemoryStore: MemorySt
18/05/28 16:07:02 INFO FileSourceStrategy: Pruning directo
18/05/28 16:07:02 INFO FileSourceStrategy: Post-Scan Filt
18/05/28 16:07:02 INFO FileSourceStrategy: Outout Data Sch

```

